

# Ruslan Sakevych

github.com/lionell

xlionell@gmail.com

1-650-732-8050

## EDUCATION

---

- **University of Washington** Seattle, USA  
*M.S. in Computer Science & Engineering; GPA: 3.83* 2022 - 2025
- **Taras Shevchenko National University** Kyiv, Ukraine  
*M.S. in Computer Science; GPA: 3.94* 2018 - 2020
- **Taras Shevchenko National University** Kyiv, Ukraine  
*B.S. in Computer Science; GPA: 3.98* 2014 - 2018

## EXPERIENCE

---

- **Google** Seattle, WA  
*Software Engineer* May 2021 - Present
  - Lead development of the critical service responsible for microarchitecture selection for executables.
  - Prototyped Reinforcement Learning based solution for the architecture selection problem.
  - Performed multiple A/B experiments that identified problematic non-cross platform code in the codebase.
  - Worked on the neural embeddings for the Google's build graph to approximate distance oracles problem.
  - Prototyped time-series forecasting (SARIMAX based) solution to mitigate the issue below.
  - Anticipated poor system accuracy due to volatility of core metrics used for control.
  - Developed online log clustering algorithm that outperformed state-of-the-art LogParser at task.
  - Performed data modeling around A/B experimentation framework that GCP uses to track VM experiments.
  - Incorporated various data sources (Linux kernel subsystems, hypervisor, drivers, etc) into VM pipelines.
- **Facebook** Seattle, WA  
*Software Engineer* Oct 2019 - May 2021
  - Joined Test Infrastructure team which later became Continuous Integration Analytics team (CIA).
  - Designed and implemented real-time service for Probabilistic Flakiness Score computation via Bayesian inference.
  - Scaled PFS service CPU utilization 10x, latency 20x, number of replicas 7.5x via profiling/caching/parallelism.
  - Details: <https://engineering.fb.com/2020/12/10/developer-tools/probabilistic-flakiness>
  - Replaced statistical model with stochastic ML one for sizing(memory-wise) CI workflows.
  - New Job Sizer model reduced memory utilization by 30% and out-of-memory rate by 17%.
  - Achieved another 12% improvement in MSLE by doing feature engineering for Job Sizer model.
  - Prototyped an ML based model for the CI workflow ETA estimation.
  - It reduced MSLE by 50% and allows to control confidence levels of the prediction (Monte-Carlo dropout layer).
- **Google** Sunnyvale, CA  
*Software Engineering Intern* May 2019 - Aug 2019
  - Designed and implemented an automatic build memory regression finder that operated at Google scale.
  - Created a developer dashboard for troubleshooting memory regression issues.
  - Was able to automatically(via dashboard) pinpoint culprit changes for recent major regressions.
- **Google** Sunnyvale, CA  
*Software Engineering Intern* Aug 2018 - Nov 2018
  - Migrated old ML pipeline onto Tensorflow-backed framework TFX. Experience with data processing pipelines.
  - Implemented parallel N-ary search algorithm. Sped culprit finder up 16x times (on millions of changes).
  - Designed more sophisticated and robust parallel batching algorithm. Reduced tail request latency 3x times.
  - Mined build graph of the whole Google using MapReduce. Did an attack on dependency set similarity problem.

- **Facebook** London, UK  
*Software Engineer Intern* Jan 2018 - Mar 2018
  - Rearchitected Hack parser to be reactive, allowing parsing to be inlined with the computation of the result.
  - 25% parse time reduction for the Hack type-checker (using most of the file contents) on the full-codebase.
  - Up to 50% speed up for tools that use less information(facts extraction) on hundreds of thousands of files.
  - Developed a toolset to analyze and remove unnecessary build dependencies, resulting in 2x speed up.
- **Microsoft** Redmond, WA  
*Software Developer Intern* Jul 2017 - Oct 2017
  - Engineered a new workflow to automate raw telemetry data aggregation and transformation.
  - System monitors execution of user-defined query and publishes results back to data warehouse.
  - Used for intermediate metrics aggregation to reduce data volumes and speed up queries.
- **Google** Sunnyvale, CA  
*Software Engineering Intern* Apr 2017 - Jul 2017
  - Research on build/test time prediction. Performed data analysis, model evaluation and feature engineering.
  - Created tools for ML models debugging/visualization and core service efficiency evaluation.
  - Investigated and mitigated incidents in complex build infrastructure at Google scale.
- **Google** Mountain View, CA  
*Software Engineering Intern* May 2016 - Aug 2016
  - Engineered a service that clusterizes build targets to reduce overall resources usage.
  - Performed evaluation of different batching strategies: memory, run-time optimization.
  - Trained ML models to predict build memory usage and avoid out of memory errors.

## SKILLS

---

- **ML:** DNN, NLP, Stochastic models, Bayesian inference, Tensorflow, TFX, PyTorch, PyTorch Lightning
- **Systems:** Distributed systems, Build/Test infrastructure, CI internals, Build systems, Programming languages
- **Languages:** Python, Java, C/C++, Go, PHP/Hack, OCaml, C#, SQL, Bash, TypeScript, JavaScript
- **Containers:** Docker, Docker Swarm, Kubernetes, RunC, Linux namespaces, Google's Borg, Facebook's Tupperware
- **Data processing:** MapReduce, Apache Beam, Hive, Hadoop, Google Flume, Facebook DataSwarm, Google Cloud SQL, Spark SQL, Microsoft Kusto
- **Misc.:** Algorithms & Data structures, Linux internals, Capture the Flag competitions, Competitive programming

## PROJECTS

---

- **PARCS:** Communication Sequential Processes (CSP) like approach for language agnostic distributed computing.
- **PARCS autodiscovery:** LAN service autodiscovery for nodes in PARCS cluster based on UDP broadcasting.
- **Smart Pacmans:** Training NN using genetic algorithms (no gradient). See <https://lionell.github.io/smart-pacmans>
- **Resolution Theorem Prover:** First-order logic theorem prover that utilizes sequential method under the hood.
- **Huffman+RLE:** Small (1% on average) optimization for Huffman algorithm followed by run-length-encoding.
- **Aqua Lang:** Imperative data processing language that uses raw relational algebra operators. Opposite to SQL (declarative).
- **Parallel PageRank:** Implemented and compared PageRank on top of MPI and OpenMP when applied to Wikipedia page graph.
- **Pollard-Rho:** Parallel implementation of Pollard-Rho algorithm in Go. **Extra:** Ethereum smart-contract impl.